



Search (II)

Mingsheng Long

Tsinghua University

Spring 2023

Outline

- **Constraint Satisfaction Problems**


- Backtracking Search
- Dynamic Ordering
- Arc Consistency
- Problem Structure

- Local Search

- Hill Climbing
- Simulated Annealing
- Local Beam Search
- Genetic Algorithm

Search Problems

- Planning: sequences of actions
 - The path to the goal is important
 - Paths have various costs and depths
 - Heuristics give problem-specific guidance




- Identification: assignments to variables
 - The goal itself is important, not the path
 - All paths generally at the same depth
 - Constraint Satisfaction Problems (CSPs)
are a special class of identification problems



Constraint Satisfaction Problems

- Standard search problems

- State is a **black box**: arbitrary data structure
- Goal test can be any function over states
- Successor function can also be anything



- Constraint satisfaction problems (CSPs)

- A special subset of search problems
- State is **structured**:
 - An assignment of variables X_i with values from a domain D_i
- Goal test is **a set of constraints** specifying allowable combinations of values for subsets of variables



Constraint Satisfaction Problems

- Constraint satisfaction problem $P = (X, D, C)$
 - Variables: $X = \{X_1, \dots, X_n\}$
 - Domains: $D = \{D_1, \dots, D_n\}$
 - Each domain $D_i = \{v_1, \dots, v_k\}$ for variable X_i
 - Constraints: C , specifying allowable combinations of values
- An assignment of values to some or all of the variables:
$$\{X_i = v_i, X_j = v_j, \dots\}$$
- A complete assignment is one in which every variable is assigned
- A solution to a CSP is a consistent and complete assignment

Example: Map Coloring

- Variables:

WA, NT, Q, NSW, V, SA, T

- Domains:


$$D = \{\text{R}, \text{G}, \text{B}\}$$

- Constraints: adjacent regions must have different colors, e.g.:

$$\text{WA} \neq \text{NT}, \dots$$

- Solutions: assignments satisfying all constraints, e.g.:

$$\begin{aligned} &\{\text{WA} = \text{R}, \text{NT} = \text{G}, \text{Q} = \text{R}, \\ &\text{NSW} = \text{G}, \text{V} = \text{R}, \text{SA} = \text{B}, \text{T} = \text{G}\} \end{aligned}$$



Example: N-Queens



- **Variables:** Q_k , the row coordinate for each column
- **Domains:** $\{1, 2, 3, \dots, N\}$
- **Constraints:**

- Implicit:

$$\forall i, j, \text{non-threatening}(Q_i, Q_j)$$

- Explicit:

$$(Q_1, Q_2) \in \{(1,3), (1,4), \dots\}, \dots$$



- **Solutions:** complete assignments for each column, e.g.:

$$Q_1 = 3, Q_2 = 1, Q_3 = 4, Q_4 = 2$$

Example: Sudoku

- Constraints:

$\text{Alldiff}(A1, A2, A3, A4, A5, A6, A7, A8, A9)$

$\text{Alldiff}(B1, B2, B3, B4, B5, B6, B7, B8, B9)$

...

$\text{Alldiff}(A1, B1, C1, D1, E1, F1, G1, H1, I1)$

$\text{Alldiff}(A2, B2, C2, D2, E2, F2, G2, H2, I2)$

...

Beyond binary constraints:

A **global constraint** is one involving an arbitrary number of variables (but not necessarily all the variables)



	1	2	3	4	5	6	7	8	9
A			3		2		6		
B	9			3		5			1
C			1	8		6	4		
D			8	1		2	9		
E	7								8
F			6	7		8	2		
G			2	6		9	5		
H	8			2		3			9
I			5		1		3		

	1	2	3	4	5	6	7	8	9
A	4	8	3	9	2	1	6	5	7
B	9	6	7	3	4	5	8	2	1
C	2	5	1	8	7	6	4	9	3
D	5	4	8	1	3	2	9	7	6
E	7	2	9	5	6	4	1	3	8
F	1	3	6	7	9	8	2	4	5
G	3	7	2	6	8	9	5	1	4
H	8	1	4	2	5	3	7	6	9
I	6	9	5	4	1	7	3	8	2

Example: The Waltz Algorithm

- Interpreting line drawings of solid polyhedra

- One of the earliest example of an AI computation posed as a CSP




What kind of intersection?
Concave or convex?

Variables: intersections


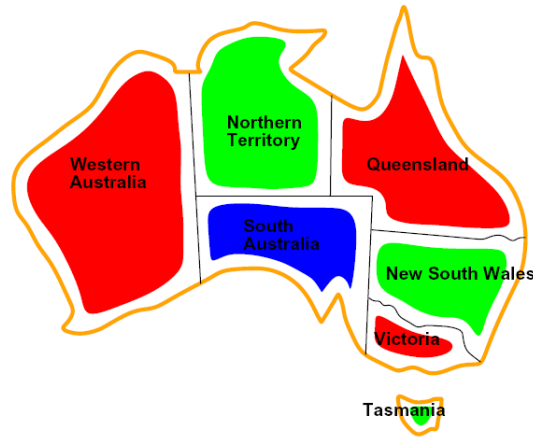
Constraints: adjacent intersections

Solutions: physically realizable 3D interpretations



Constraint Graphs


- Binary CSP: each constraint relates (at most) two variables
- Binary constraint graph: nodes are variables, arcs show constraints



- General-purpose CSP algorithms:
 - We use the graph structure to speed up search (more later)
 - E.g., Tasmania (T) is an independent subproblem

Applications of CSPs

- Assignment problems: e.g., who teaches what class
 - Timetabling problems: e.g., which class is offered when and where
 - Transportation scheduling
 - Factory scheduling
 - Hardware configuration
 - Circuit layout
 - Fault diagnosis
 - And lots more...
- Many real-world problems involve **real-valued** variables and difficult...



Outline

- **Constraint Satisfaction Problems**
 - **Backtracking Search**
 - Dynamic Ordering
 - Arc Consistency
 - Problem Structure
- Local Search
 - Hill Climbing
 - Simulated Annealing
 - Local Beam Search
 - Genetic Algorithm

Standard Search Formulation



Standard search formulation of CSPs:

- **States:** the values assigned so far (**partial assignments**)
- **Initial state:** the empty assignment, $\{\}$
- **Successor function:** **assign a value** to an unassigned variable
- **Goal test:** complete and satisfy all constraints

- Standard search methods: DFS, BFS, ...
- What problems does naive search have?




Map Coloring: Depth First Search



DFS:
Duplicate
states and
paths

	WA	NT	Q	NSW	V	SA
Queue #1						
#2						
#3						
#4						
#5						
#6						



Backtracking Search

- Idea I: One variable at a time
 - Variable assignments are commutative, so **fix ordering**
 - I.e., $[WA = R \text{ then } NT = G]$ same as $[NT = G \text{ then } WA = R]$
- Idea II: Check constraints as you go
 - **Incremental goal test**: consider only values which do not conflict with previous assignments
 - Might have to do some computation to check the constraints

Backtracking Search = DFS + **variable-ordering** + **fail-on-violation**

Backtracking Search

- Backtracking Search = DFS + 变量排序 variable-ordering + fail-on-violation

function BACKTRACK(*assignment*, *Domains*) **returns** a solution, or *failure*

if *assignment* is complete **then return** *assignment* ← General Search checks consistency on full assignment

choose unassigned variable X_i

order domain values $Domains_i$ of chosen X_i 对其他值排序

Variable-ordering

for each *value* **in** that order **do**:

if *value* is consistent with *assignment* **then** ←

Fail-on-violation

add { $X_i = value$ } to *assignment*

result = BACKTRACK(*assignment*, *Domains*)

Backtracking Search checks consistency at each assignment


if *result* \neq *failure* **then return** *result*

remove { $X_i = value$ } from *assignment* ←

Backtracking (回溯)

return *failure*

Backtracking Example



How to Improve Backtracking?

- General-purpose ideas give **huge gains** in speed

- Heuristics

- Ordering


- Which **variable** should be assigned next?
- In what order should its **values** be tried?

- Filtering

- Can we detect inevitable failure early?

- Structure

- Can we exploit the problem structure?



Outline

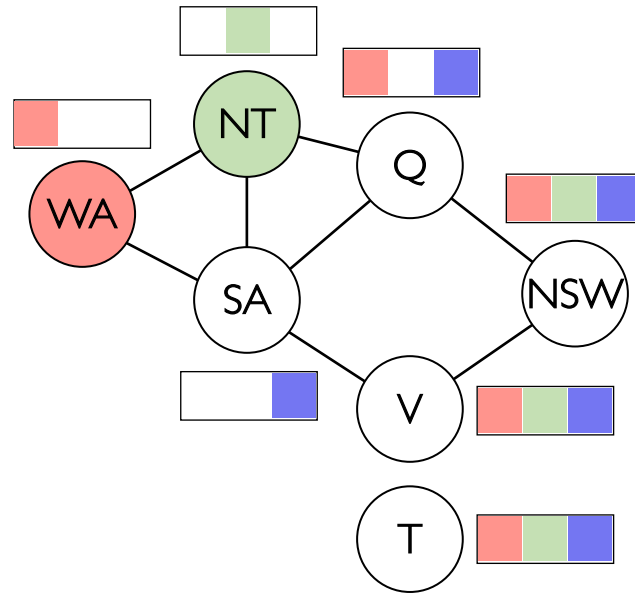
- **Constraint Satisfaction Problems**

- Backtracking Search
- **Dynamic Ordering**
- Arc Consistency
- Problem Structure

- Local Search Methods

- Hill Climbing
- Simulated Annealing
- Local Beam Search
- Genetic Algorithm

Choose an Unassigned Variable



- Which variable to assign next?

Smaller branching factors


Key idea: most constrained variable

Choose variable that has the fewest consistent values.

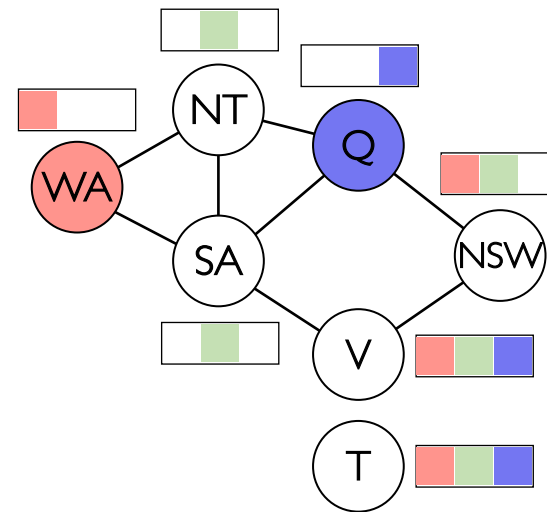
- This example: SA (has only one value)

Order Values of a Selected Variable

- What values to try for Q?



$2 + 2 + 2 = 6$ consistent values



$1 + 1 + 2 = 4$ consistent values


Key idea: least constrained value

Order values of selected X_i by decreasing number of consistent values of neighboring variables.

Resolve Conflicting Heuristics


- Most constrained variable (MCV):

- Must assign **every** variable
- If going to fail, fail early → more pruning
- For some problem, improve by **1000x**



- Least constrained value (LCV):

- Need to choose **some** value
- Choosing value most likely to lead to solution



- Combining these ordering ideas makes **1000 queens** feasible

Outline

- **Constraint Satisfaction Problems**

- Backtracking Search
- Dynamic Ordering
- **Arc Consistency**
- Problem Structure

- Local Search

- Hill Climbing
- Simulated Annealing
- Local Beam Search
- Genetic Algorithm

Interleave Search and Inference

- **Filtering**: Keep track of domains for unassigned variables and **infer new domain reductions**

function **BACKTRACK**(*assignment*, *Domains*) **returns** a solution, or *failure*

if *assignment* is complete **then return** *assignment*

choose unassigned variable X_i

order domain values $Domains_i$ of chosen X_i

for each *value* **in** that order **do**:

if *value* is consistent with *assignment* **then**

add { $X_i = \text{value}$ } to *assignment*

$Domains_i = \{\text{value}\}$

if **FILTER**(X_i , *Domains*) \neq *failure* **then**

result = **BACKTRACK**(*assignment*, *Domains*)


if *result* \neq *failure* **then return** *result*

remove { $X_i = \text{value}$ } from *assignment* and restore *Domains*

return *failure*

Filtering (Forward checking or AC-3) removes bad values from domains.
If one of the domains is reduced to empty, prune.


Filtering: Forward Checking



Filtering: Forward Checking

- Forward checking (one-step lookahead)

- After assigning a variable X_i , eliminate inconsistent values from the domains of X_i 's neighbors
- If any domain becomes empty, don't recurse.
- When unassign X_i , restore neighbors' domains.



Arc Consistency

- An arc $X_i \rightarrow X_j$ is **consistent** iff for **every** $x_i \in \text{Domains}_i$, there exists $x_j \in \text{Domains}_j$ which can be assigned **without violating a constraint**
- **Enforce arc consistency**: Remove values from Domains_i to make arc consistent

function **EnforceArcConsistency**($\text{Domains}, X_i, X_j$) **returns** *true* iff succeeds

removed = false

for each x **in** Domains_i **do**

if no value y in Domains_j allows (x, y) to satisfy the constraint between X_i and X_j

then delete x from Domains_i ; *removed = true*


return *removed*

Remember: Delete from Domains_i

- **Forward checking**: Enforcing consistency of arcs pointing to each new assignment

Filtering: Constraint Propagation

- Forward checking doesn't provide early detection for **all failures**
- **Constraint propagation**: reasoning from constraint to constraint
 - Algorithms: **AC-3** (the most popular), PC-2, etc.
 - Detects failure **earlier** than forward checking



Constraint propagation:
further remove **Blue** for Q

AC-3

- **AC-3**: repeatedly enforce arc consistency on constraint chains
- **Important**: If X loses a value, neighbors of X need to be **rechecked**

Called when the domain of X_t is reduced.

```
function AC-3( $X_t$ , Domains) returns false if an inconsistency is found and true otherwise
  initialize queue with all arcs ( $X_s, X_t$ ) for  $X_s$  in Neighbours( $X_t$ )
  while queue is not empty do
    ( $X_i, X_j$ )  $\leftarrow$  RemoveFirst(queue)
    if EnforceArcConsistency(Domains,  $X_i, X_j$ ) then The domain of  $X_i$  is reduced.
      if size of Domains $i$  = 0 then return false
      for each  $X_k$  in Neighbours( $X_i$ ) do
        add ( $X_k, X_i$ ) to queue
  return true
```

Constraint propagation:



AC-3 Properties


- Assume a CSP with:
 - n variables each with domain size at most d , c binary constraints
- The complexity of AC-3: $O(c \cdot d \cdot d^2)$

For each arc $X_k \rightarrow X_i$

Runtime for enforcing arc consistency

Inserted in the queue up to d times because X_i has at most d values to delete

- AC-3 isn't always effective:
 - No consistent assignments, but AC-3 doesn't detect it...
 - **Intuition**: if we look **locally** at the graph, nothing blatantly wrong.



Outline

- **Constraint Satisfaction Problems**


- Backtracking Search
- Dynamic Ordering
- Arc Consistency
- **Problem Structure**

- Local Search

- Hill Climbing
- Simulated Annealing
- Local Beam Search
- Genetic Algorithm

Problem Structure

- Extreme case: **independent subproblems**
 - Example: Tasmania and mainland do not interact
- Independent subproblems are identifiable as **connected components** of constraint graph




Example: Suppose a graph of n variables can be broken into subproblems of only c variables taking d values:

- Worst-case solution cost is $O(d^c n/c)$, linear in n
- E.g., $n = 80, d = 2, c = 20$, 10 million nodes/sec
- Without the decomposition: $2^{80} = 4$ billion years
- With the decomposition: $4 \cdot 2^{20} = 0.4$ seconds

Tree-Structured CSPs

- **Theorem:** if the constraint graph has **no loops**, the CSP can be solved in linear time: $O(nd^2)$
 - Compare to general CSPs, where worst-case time is $O(d^n)$




- This property also applies to **probabilistic reasoning on graphs**
 - An example of the relation between syntactic restrictions and the complexity of reasoning

Tree-Structured CSPs


- Algorithm for tree-structured CSPs:

Runtime: $O(nd^2)$

- **Ordering**: Choose a root variable, order variables so that parents precede children



Topological
Sort



- Remove backward:

- Assign forward:

```

for  $j = n$  to 2 do
    EnforceArcConsistency(Parent( $X_j$ ),  $X_j$ )
    if it cannot be made consistent then return failure

for  $i = 1$  to  $n$  do
     $assignment[X_i] \leftarrow$  any consistent value from  $D_i$ 
    if there is no consistent value then return failure
    
```

Outline


- Constraint Satisfaction Problems
 - Backtracking Search
 - Dynamic Ordering
 - Arc Consistency
 - Problem Structure
- **Local Search**
 - Hill Climbing
 - Simulated Annealing
 - Local Beam Search
 - Genetic Algorithm

Local Search


- Solving CSPs is often NP-hard
 - Huge search space: exponential in the number of variables
 - All have cost $O((\max_i |\text{Domains}_i|)^n)$, only useful for constants
- Alternative: Local search
 - Generally much faster and memory efficient (a constant amount)
 - But incomplete and suboptimal
- Useful method in practice
 - Best available method for many constraint satisfaction and constraint optimization problems
 - Online processing, e.g. fast flight reschedule due to weather change

Solving CSPs: Local Search


- **Tree search** keeps unexplored alternatives on the frontier
 - Ensures completeness and optimality
 - Extends **partial assignments**
- **Local search** improves **a single option** until you can't make it better
- Apply local search to CSPs:
 - Take a **complete assignment** with **unsatisfied** constraints
 - **Reassign** variable values until solved
 - No frontier
 - Live on the edge




Applications of Local Search



Probabilistic Reasoning



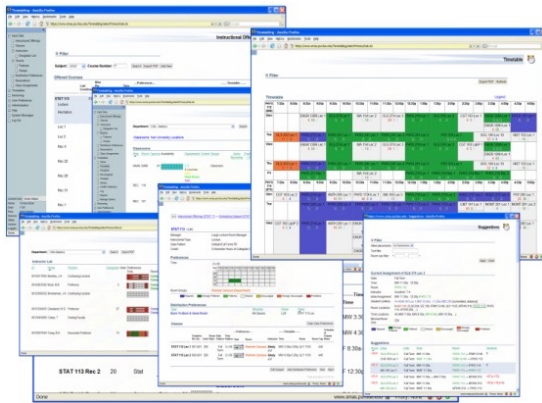
RNA structure design




Propositional satisfiability (SAT)



Scheduling of Hubble
Space Telescope:
1 week \rightarrow 10 seconds



University Timetabling



Protein Folding

Outline

- Constraint Satisfaction Problems
 - Backtracking Search
 - Dynamic Ordering
 - Arc Consistency
 - Problem Structure
- **Local Search**
 - **Hill Climbing**
 - Simulated Annealing
 - Local Beam Search
 - Genetic Algorithm

Hill Climbing


- Also called **greedy local search**
- Simple and general idea:
 - Start wherever
 - **Repeat:** move to the **best neighboring state**
 - If no neighbors better than current, quit




```
function HILL-CLIMBING(problem) returns a state that is a local maximum  
  current  $\leftarrow$  problem.INITIAL  
  while true do  
    neighbor  $\leftarrow$  a highest-valued successor state of current  
    if VALUE(neighbor)  $\leq$  VALUE(current) then return current  
    current  $\leftarrow$  neighbor
```


Example: N-Queen


- At each stage, a queen is chosen for **reassignment in its column**.
- **The number of conflicts** (in this case, the number of attacking queens) is shown in each square.



Random
initialization




Reassignment for
a column



Reassignment for
another. Done

Hill Climbing Diagram

- Problems:
 - **Local optima**: a peak that is higher than each of its neighbors
 - **Get stuck** with nowhere else to go
 - **Plateaus**: a flat local maximum or a shoulder
 - **Get lost**: unable to determine in which direction it should step




Local Optima

- Current cost: $h = 1$
- No single move can improve on this solution
 - In fact, every single move only makes things worse ($h \geq 2$)

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18

$h = 17$


5 steps



$h = 1$

How to Improve Hill Climbing?

- **Stochastic hill climbing**: chooses **at random** from uphill moves
 - The probability of selection can **vary with the steepness**
- **Random-restart hill climbing**: a series of hill-climbing searches from **randomly generated initial states**, until a goal is found
 - Complete with probability approaching 1




For 8-queens,
each hill-climbing search has
a probability p of success
 $p \approx 0.14$, meaning roughly 7
iterations for global optima

Outline

- Constraint Satisfaction Problems
 - Backtracking Search
 - Dynamic Ordering
 - Arc Consistency
 - Problem Structure
- **Local Search**
 - Hill Climbing
 - **Simulated Annealing**
 - Local Beam Search
 - Genetic Algorithm

Simulated Annealing

- **Simulated Annealing**: physics inspired twist on random walk




- **Basic idea:**

- Allow "bad" moves occasionally
- With high temperature, more bad moves allowed, shake the system out of its local minimum
- Gradually reduce temperature



Imagine a ping-pong



Simulated Annealing

function SIMULATED-ANNEALING(*problem*, *schedule*) **returns** a solution state

current \leftarrow *problem*.INITIAL

for $t = 1$ **to** ∞ **do**

$T \leftarrow$ *schedule*(t)

if $T = 0$ **then return** *current*

next \leftarrow a randomly selected successor of *current*

$\Delta E \leftarrow$ VALUE(*current*) – VALUE(*next*)

if $\Delta E > 0$ **then** *current* \leftarrow *next*


else *current* \leftarrow *next* only with probability $e^{\Delta E/T}$

Acceptance

Probability ($\Delta E < 0$):

$$e^{\Delta E/T}$$

- **Cooling schedule**: a gradual reduction from a high value to 0
 - **Exponential cooling** often works best, typically at a rate of **0.7~0.9**
 - Complete with probability approaching **1**




Outline


- Constraint Satisfaction Problems
 - Backtracking Search
 - Dynamic Ordering
 - Arc Consistency
 - Problem Structure
- **Local Search**
 - Hill Climbing
 - Simulated Annealing
 - **Local Beam Search**
 - Genetic Algorithm

Local Beam Search


- Basic idea: greedily keep k states at all times
 - Begins with k randomly generated states
 - For each iteration
 - Generate all successors from k current states
 - Choose best k of these to be the new current states



Greedy Search



Random-restart
Greedy Search




Local Beam Search ($k = 2$)

Information is shared
among k search threads.


Beam Search vs. Local Beam Search

- **Beam Search**: a path-based algorithm
 - Not guaranteed to find optimal assignment
- Running time: $O(n(kb) \log(kb))$ with branching factor b

Beam
Search
($k = 3$)




Local Beam
Search
($k = 3$)



Stochastic Beam Search

- Local beam search can suffer from **lack of diversity** among k states
 - They can quickly become concentrated in a small region of the state space.
- **Stochastic beam search**
 - Chooses k from the the pool of candidate successors **at random**
 - **Probability of chosen**: an increasing function of its value




Outline


- Constraint Satisfaction Problems
 - Backtracking Search
 - Dynamic Ordering
 - Arc Consistency
 - Problem Structure
- **Local Search**
 - Hill Climbing
 - Simulated Annealing
 - Local Beam Search
 - **Genetic Algorithm**

Genetic Algorithm (GA)



- Basic Idea:
 - A variant of **stochastic beam search**
 - **Successor states** are generated by **combining two parent states** rather than by modifying a single state
 - Hill Climbing + Stochastic Exploration + Parallel Communication
- Example: 8-queens



+

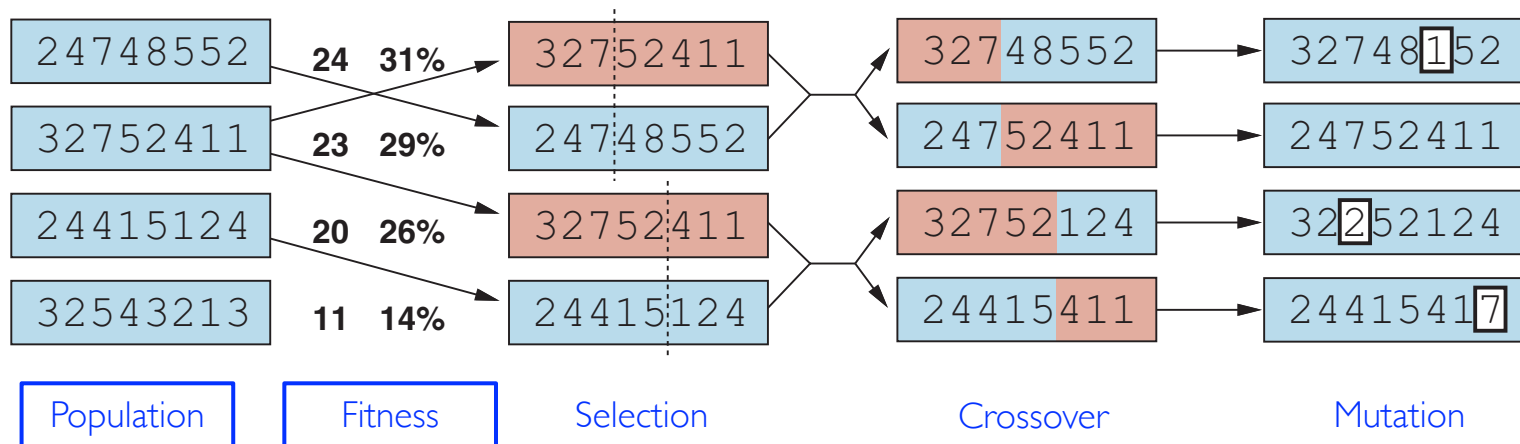


=



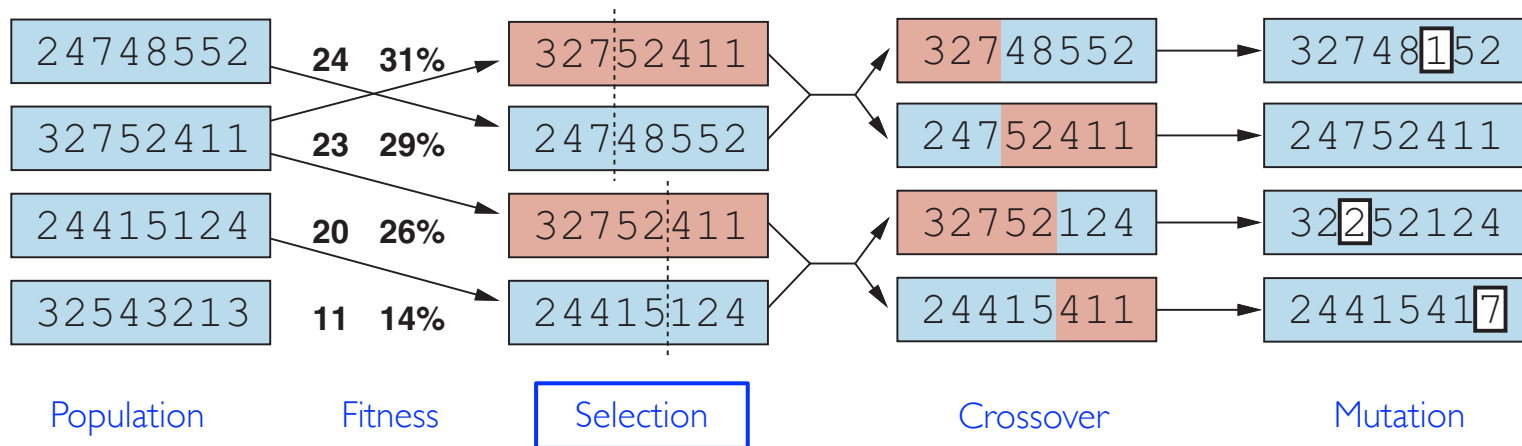
Genetic Algorithm: Population

- **Population:** initially a set of k randomly generated states
- **Individual (state):** represented as a **chromosomes** over **a finite alphabet**
 - 8-queens problem: $Q_1 Q_2 \cdots Q_8$
- **Fitness function:** the objective function
 - 8-queens problem: the number of nonattacking pairs of queens



Genetic Algorithm: Selection

- **Selection**: pairs are selected at random for **reproduction**
 - The probability of being chosen is directly **proportional to the fitness score**
 - There are also many variants of this selection rule
 - **Knowledge engineering** is important: should select **modular parts**



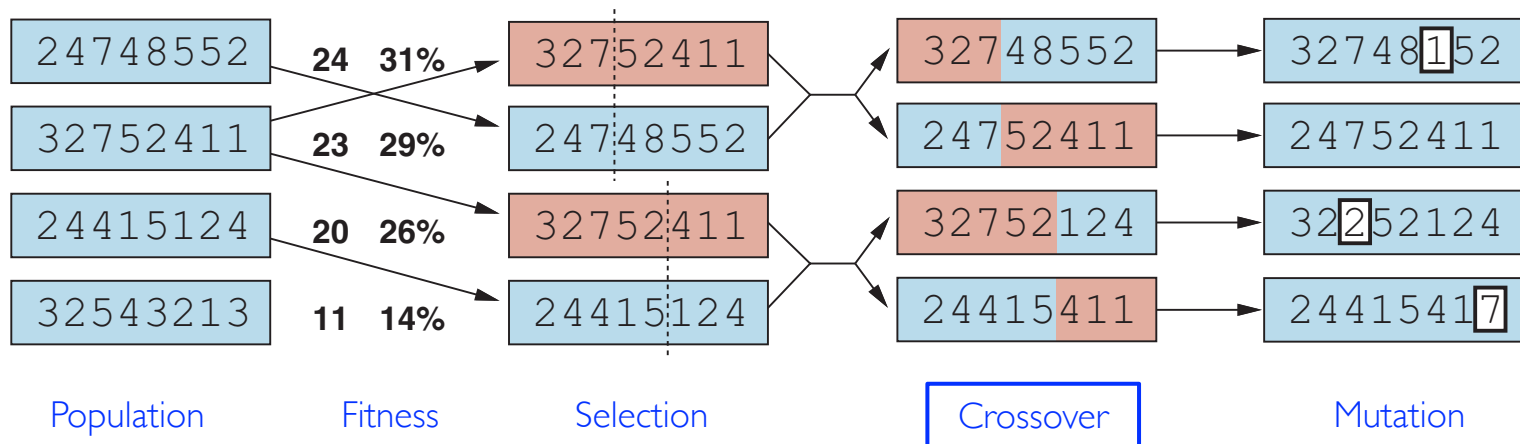
Genetic Algorithm: Crossover

- **Crossover**: offspring are created by crossing over the parent strings
 - **Crossover point**: chosen randomly from the positions in the string

Intuition: Like Simulated Annealing, crossover

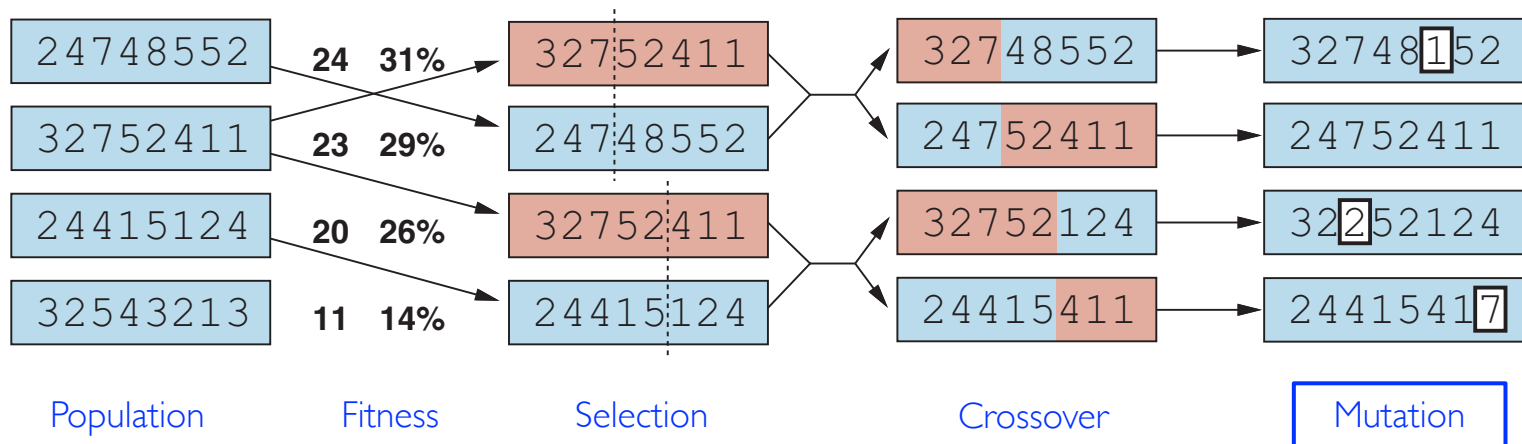
- **Early**: takes **large** steps with **diverse** population
- **Later**: takes **smaller** steps with **similar** individuals

Simulated Annealing



Genetic Algorithm: Mutation

- **Mutation**: each location is subject to **random mutation** with a **small independent probability**
 - Mutation increases the diversity of individuals in the population
 - Good mutation with higher fitness score will gain more popularity
- The fittest survive by natural selection. -- Darwinism



Genetic Algorithm

function GENETIC-ALGORITHM(*population*, *fitness*) **returns** an individual
repeat

weights \leftarrow WEIGHTED-BY(*population*, *fitness*)

population2 \leftarrow empty list

for *i* = 1 **to** SIZE(*population*) **do**

Selection *parent1*, *parent2* \leftarrow WEIGHTED-RANDOM-CHOICES(*population*, *weights*, 2)

Crossover *child* \leftarrow REPRODUCE(*parent1*, *parent2*)

Mutation **if** (small random probability) **then** *child* \leftarrow MUTATE(*child*)

 add *child* to *population2*

population \leftarrow *population2*

until some individual is fit enough, or enough time has elapsed

return the best individual in *population*, according to *fitness*


function REPRODUCE(*parent1*, *parent2*) **returns** an individual

n \leftarrow LENGTH(*parent1*)


c \leftarrow random number from 1 to *n*

return APPEND(SUBSTRING(*parent1*, 1, *c*), SUBSTRING(*parent2*, *c* + 1, *n*))

Crossover details




Small Picture of Local Search




Random-restart
Hill Climbing

Random Restart



Hill Climbing
(Greedy Local Search)

Keeping k states




Local Beam Search

Randomization

Stochastic Hill
Climbing

Allowing
bad moves

Simulated Annealing




Randomization

Stochastic Beam Search

Evolutionism

Genetic
Algorithm
(GA)



Thank You

Questions?

Mingsheng Long

mingsheng@tsinghua.edu.cn

<http://ise.thss.tsinghua.edu.cn/~mlong>

答疑：东主楼11区413室

[Some slides adapted from Dan Klein and Pieter Abbeel]