



Reasoning (I)

Mingsheng Long

Tsinghua University

Spring 2023

Outline

- **Probabilistic Reasoning**
- Graphical Models
 - Bayesian Networks
 - Markov Random Fields
- Conditional Independence
 - D-separation
 - Markov Properties
- Inference in Graphical Models
 - Variable Elimination
 - Belief Propagation

Modeling and Inference

- How does a weather forecast predict the probability of Sunshine (S) or Rain (R) given the observation of Temperature (T) and Date (D)?

- **Modeling:**

Find the **joint** distribution
 $p(S, R, T, D)$

- **Inference:**

- **Condition** on evidence (Temperature and Date):

$$T = 36.5^{\circ}\text{C}, D = 2021/7/31$$

- Interested in certain variable set (**query**), such as rain ($\{R\}$):

$$p(R \mid T = 36.5^{\circ}\text{C}, D = 2021/7/31)$$

(S is marginalized out)

Probabilistic Reasoning

- **Probabilistic Reasoning = Modeling + Inference**
 - $X = \{x_1, \dots, x_D\}$ is a set of D random variables.
 - Query set R and condition set C are subsets of X .
- **Modeling**: How to specify a joint distribution $p(x_1, \dots, x_D)$ compactly?
 - Bayesian networks
 - Markov random fields
- **Inference**: How to compute $p(R \mid C)$ efficiently?
 - Elimination methods (变量消除法)——lec9
 - Latent variable models (隐变量模型)——lec10
 - Variational methods (变分方法)——lec11
 - Sampling methods (采样方法)——lec12

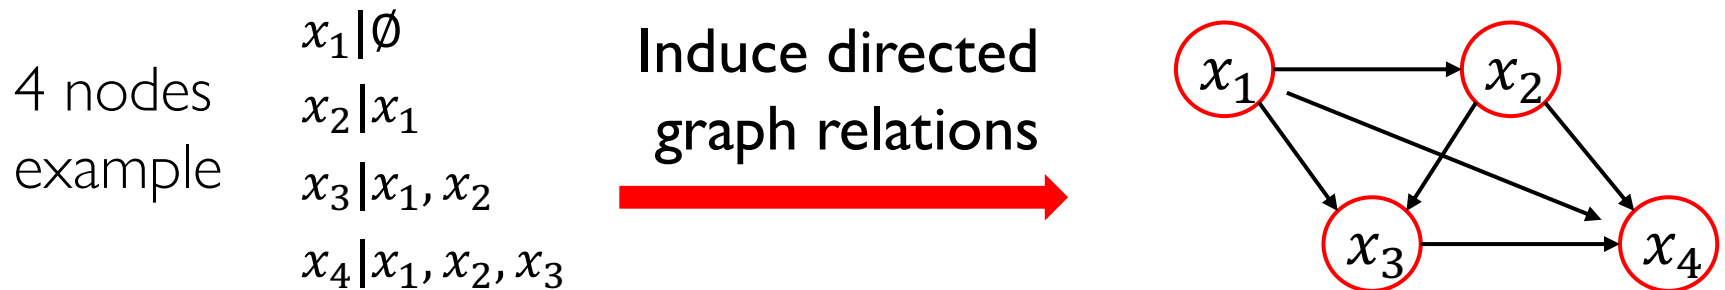
Outline

- Probabilistic Reasoning
- **Graphical Models**
 - **Bayesian Networks**
 - Markov Random Fields
- Conditional Independence
 - D-separation
 - Markov Properties
- Inference in Graphical Models
 - Variable Elimination
 - Belief Propagation

Bayesian Network

- Product rule of probability implies joint distribution of K variables:

$$p(x_1, \dots, x_K) = p(x_K | x_1, \dots, x_{K-1}) \cdots p(x_2 | x_1) p(x_1).$$



- A **Bayesian network** is a **directed acyclic graph (DAG)** that specifies a joint distribution as a product of local conditional distributions, one for each node:

$$p(x_1, \dots, x_K) = \prod_{s=1}^K p(x_s | \mathbf{x}_{\Gamma(s)})$$

where $\Gamma(s)$ denotes the set of parents of x_s .

Example: The Alarm Network

Q: Earthquakes (E) and burglaries (B) are independent events that will cause an alarm (A) to go off. Suppose you **hear an alarm**. How does **hearing the earthquake report** (R) change your beliefs about burglary?



Burglaries ↘



↙ Earthquakes ↘

Earthquakes
Report



Alarm

**Hearing
the alarm**



Radio

Example: The Alarm Network

- Choosing an **ordering**:

$$p(A, R, E, B) = p(A|R, E, B)p(R|E, B)p(E|B)p(B)$$

- Assumptions:**

- The alarm is not influenced by the radio report.

$$p(A|R, E, B) = p(A|E, B)$$




- The radio does not broadcast any burglary activity.

$$p(R|E, B) = p(R|E)$$

- Burglaries don't **cause** earthquakes.

$$p(E|B) = p(E)$$


- Modeling:** $p(A, R, E, B) = p(A|E, B)p(R|E)p(E)p(B)$



Example: The Alarm Network

- Parameters:

- $p(B = 1) = p(E = 1) = 0.01$
- $B \vee E \Rightarrow (A = 1)$ almost surely
- $E \Rightarrow R$ almost surely



- Inference:

- Probability of Burglary **only given Alarm**:

$$p(B = 1|A = 1) = \frac{\sum_{E,R} p(B = 1, E, A = 1, R)}{\sum_{B,E,R} p(B, E, A = 1, R)} \approx 0.5$$

Earthquakes or Burglaries
almost half and half



- Probability of Burglary **given Alarm and Radio**:

$$p(B = 1|A = 1, R = 1) = 0.01$$




Nearly sure the
Earthquakes

Explaining away


Example: Medical Diagnosis

Explainable

Medical illness diagnoses



ICU monitoring



Example: Polynomial Regression

- The polynomial regression model, jointly showing also a new input value \hat{x} together with the corresponding model prediction \hat{t} .


给定 x_{hat} , t_{hat} 发生的概率

$$p(\hat{t}, \mathbf{t}_{1:N}, \mathbf{w} | \hat{x}, \mathbf{x}_{1:N}, \alpha, \sigma^2) = \left[\prod_{n=1}^N p(t_n | x_n, \mathbf{w}, \sigma^2) \right] \times p(\mathbf{w} | \alpha) p(\hat{t} | \hat{x}, \mathbf{w}, \sigma^2)$$

通常认为是均值为 $\mathbf{w} \times x_{\text{hat}}$, 方差 σ^2 的高斯分布

parameter

hyperparameter




Outline

- Probabilistic Reasoning
- **Graphical Models**
 - Bayesian Networks
 - **Markov Random Fields**
- Conditional Independence
 - D-separation
 - Markov Properties
- Inference in Graphical Models
 - Variable Elimination
 - Belief Propagation


Are Bayesian Network Enough

- Bayesian networks (Directed acyclic graph, DAG) can hardly model the **bi-directional** or **cycling** relationships.

Adjacency between
Pixels in a Figure



Friend relation in
Social Networks



Example: Voting Preferences


- Four students A, B, C, D are voting for a group leader.
- Say that (A, B) , (B, C) , (C, D) , and (D, A) are friends, and friends tend to have similar voting preference.
- Friend relationships can be naturally represented by an undirected graph.
 - Identify dependent variables:

$$p(A, B, C, D) = \frac{1}{Z} \phi(A, B) \phi(B, C) \phi(C, D) \phi(D, A)$$

where Z is a normalizing constant to ensure p forms a distribution.

- Define the strength of interactions:

$$\phi(X, Y) = \begin{cases} 10 & \text{if } X = Y = 1 \\ 5 & \text{if } X = Y = 0 \\ 1 & \text{otherwise} \end{cases}$$




Markov Random Field

- A **Markov Random Field (MRF)** is an **undirected graph** that specifies a **joint distribution** as a product of **potential functions**:

$$p(\mathbf{x}) = \frac{1}{Z} \prod_C \psi_C(\mathbf{x}_C)$$

团

- C denotes a **clique** (fully connected subgraphs) in graph G .
 - \mathbf{x} denotes $\{x_1 \dots x_n\}$, \mathbf{x}_C denotes $\{x_i \mid x_i \in C\}$.
 - Each ψ_C is a **non-negative** function over the variables in a clique.
 - Z is the normalizing constant (partition function).
- We cannot **factorize** the potential function of a fully connected subgraph.



Energy Function

- Because we are restricted to strictly **positive potential functions**, it is convenient to express them as exponentials.

$$\psi_c(\mathbf{x}_c) = \exp\{-E(\mathbf{x}_c)\}$$


- $E(\mathbf{x}_c)$ is called an **energy function**, and so the total energy is obtained by adding the energies of each cliques.

$$\begin{aligned} p(\mathbf{x}) &= \frac{1}{Z} \prod_c \psi_c(\mathbf{x}_c) \\ &= \frac{1}{Z} \exp\left\{-\sum_c E(\mathbf{x}_c)\right\} \end{aligned}$$

Lower energy \Rightarrow Higher probability


Example: Image Denoising

Denoising



Noise-free image

$$x_i \in \{-1, +1\}$$



Observed image


$$y_i \in \{-1, +1\}$$

Flipping the sign with probability 10%

Example: Image Denoising

- Modeling image with an MRF:

- **Nodes:** pixels in both images x_i, y_i .
- **Edges:** (x_i, x_j) for each neighbor pair i, j ;
 (x_i, y_i) for each position i .



- Define the energy function (Pixels with same sign \Rightarrow Lower energy)
 - Pixels at same position x_i and y_i have strong correlation.

$$E(x_i, y_i) = -\eta x_i y_i$$

- Neighboring pixels x_i and x_j in an image are strongly correlated.

$$E(x_i, x_j) = -\beta x_i x_j$$

Example: Image Denoising

- The complete energy function

$$E(\mathbf{x}, \mathbf{y}) = -\eta \sum_i x_i y_i - \beta \sum_{\{i,j\}} x_i x_j$$

- The joint distribution

$$p(\mathbf{x}, \mathbf{y}) = \frac{1}{Z} \exp\{-E(\mathbf{x}, \mathbf{y})\}$$

- Pixels of observed image $\rightarrow \mathbf{y}$.
- **Denosing**: find noise-free image \mathbf{x} with high probability given the observation \mathbf{y} .

$$\hat{\mathbf{x}} = \operatorname{argmax}_{\mathbf{x}} p(\mathbf{x} \mid \mathbf{y})$$

Example: Image Denoising



- Algorithm (Iterated conditional modes, ICM)

1. Initialize the variables $\{x_i\}$ (e.g. by $x_i = y_i$ for all i).

2. For each pixel x_i :

Evaluate the total energy of possible states $x_i = +1$ and $x_i = -1$ keeping all other node fixed.

Set x_i to whichever state has the lower energy.

$$x_i^{\text{new}} = \underset{x_i \in \{0,1\}}{\operatorname{argmax}} p(x_i \mid \mathbf{y}, \mathbf{x}^{\text{old}})$$


-1, 0是错的

3. Repeat step 2. until converge.

Greedy Method!



Relation to Directed Graphs



$$p(B)p(E)\underbrace{p(A | B, E)}_{\text{Cannot decompensate}}p(R | E)$$

**Cannot
decompensate**

$$\phi_1 = p(B)p(E)p(A | B, E)$$

$$\phi_2 = p(R | E)$$

- Bayesian network can be transformed into MRF by
 - Linking the parents of each node;
 - Removing the direction of each arrow.
- The transformation method is not unique.

Outline

- Probabilistic Reasoning
- Graphical Models
 - Bayesian Networks
 - Markov Random Fields
- **Conditional Independence**
 - **D-separation**
 - Markov Properties
- Inference in Graphical Models
 - Variable Elimination
 - Belief Propagation

Conditional Independence

- X and Y are **independent** random variables if and only if

$$p_{XY}(x, y) = p_X(x)p_Y(y)$$

for all $x \in \mathcal{X}, y \in \mathcal{Y}$, denoted as $X \perp Y$.

- X and Y are **conditionally independent** random variables given another conditioning variable Z (also known as “**observed variable**”)


$$p_{XY|Z}(x, y | z) = p_{X|Z}(x | z) p_{Y|Z}(y | z)$$

or alternatively,

$$p_{X|YZ}(x | y, z) = p_{X|Z}(x | z)$$

for all $x \in \mathcal{X}, y \in \mathcal{Y}$ and $z \in \mathcal{Z}$, denoted as $X \perp Y | Z$.

Marginal Independence in Directed Models




- In (a), (b) and (c), A and B are **dependent** random variables.
- In (d), A and B are **independent** random variables.

Definition of
Bayesian networks

$$\begin{aligned} p(A, B) &= \sum_C p(A, B, C) = \sum_C p(A)p(B)p(C|A, B) \\ &= p(A)p(B) \sum_C p(C|A, B) = p(A)p(B) \end{aligned}$$

Conditional Independence in Directed Models



- In (a), (b) and (c), A and B are **conditionally independent** given C .

- (a) $p(A, B | C) = \frac{p(A|C) p(B|C) p(C)}{p(C)} = p(A | C) p(B | C)$

- (b) $p(A, B | C) = \frac{p(A|C) p(C|B) p(B)}{p(C)} = p(A | C) p(B | C)$ Bayes formula

- In general case of (d), A and B are **conditionally dependent** given C .


- (d) $p(A, B | C) \propto p(C | A, B) p(A) p(B)$

Marginally independent but conditionally dependent!


Conditional Dependence Criteria

- Consider an **undirected** path U between A and B , the path is **connected** given condition set \mathcal{C} if it belongs to one of these cases:

- **Non-head-to-head** nodes of U do not belong to \mathcal{C} .



- Any **head-to-head** node belongs to \mathcal{C} or it has descendant that belongs to \mathcal{C} .




D-separation


- All paths from any node in A (node set) to any node in B (node set) are **disconnected** given C .

$\Leftrightarrow A$ is said to be **D-separated** from B by C .

$\Leftrightarrow (A \perp B) \mid C$



Although c is a head-to-head node, a and b are connected due to d is observed.



D-separated!

Markov Blanket

- How to infer a node with variable x_i from the remaining variables $x_{j \neq i}$ in directed models ?

$$p(x_i \mid x_{\{j \neq i\}})$$

Worst case: Depend on all nodes of the graph

- **Simplify**: can we use part of the graph for inference?
- Find $S_1 \in x_{\{j \neq i\}}$ which contents

$$x_i \perp (x_{\{j \neq i\}} \setminus S_1) \mid S_1$$

- It means that S_1 contains all information one needs to infer x_i .

$$p(x_i \mid x_{\{j \neq i\}}) = p(x_i \mid S_1 \sqcup (x_{\{j \neq i\}} \setminus S_1)) = p(x_i \mid S_1)$$

- S_1 is called the **Markov blanket** of x_i .


Markov Boundary

- How to infer a node with variable x_i from the remaining variables $x_{j \neq i}$ in directed models ?
- A **Markov boundary** is the **minimal** Markov blanket. In a Bayesian network, it includes **parents**, **children** and the other parents of all of its children (**co-parents**).

$$p(x_i \mid \mathbf{x}_{\{j \neq i\}}) = p(x_i \mid \mathbf{x}_{pa} \cup \mathbf{x}_{ch} \cup \mathbf{x}_{co\cdot pa})$$

- $B \not\perp A \mid D$ while $C \perp A \mid B, D$

Markov Boundary 'block' every path from A to outside nodes.



Outline

- Probabilistic Reasoning
- Graphical Models
 - Bayesian Networks
 - Markov Random Fields
- **Conditional Independence**
 - D-separation
 - **Markov Properties**
- Inference in Graphical Models
 - Variable Elimination
 - Belief Propagation


Markov Properties

- What independencies can be described by an undirected MRF?
 - x, y are **dependent** if they are **connected** by a path of unobserved variables.

Global Markov Property:

If any path from \mathbf{A} to \mathbf{B} passes through \mathbf{C} , which denoted as $\text{sep}_G(\mathbf{A}, \mathbf{B} \mid \mathbf{C})$, then

$$(X_A \perp X_B) \mid X_C$$




Markov Properties

- If x 's neighbors are all observed, then x is independent of all the other variables.

Local Markov Property:

For any node x , $N(x)$ denotes neighbors of x , $N[x] = x \cup N(x)$, then

$$X_x \perp X_{V \setminus N[x]} \mid X_{N(x)}$$




- Suppose u and v are not linked, they are independent when all the other variables are observed.

Pairwise Markov Property:

If u and v are not linked, then

$$X_u \perp X_v \mid X_{V \setminus \{u,v\}}$$



Markov Random Fields

- \mathbf{X} forms a **Markov random fields** with respect to graph $G = (V, E)$ if they satisfy the Markov properties:

(a) Global Markov property:

$$\text{sep}_G(\mathbf{A}; \mathbf{B} \mid \mathbf{C}) \Rightarrow (\mathbf{X}_A \perp \mathbf{X}_B) \mid \mathbf{X}_C \text{ for all disjoint set } \mathbf{A}, \mathbf{B}, \mathbf{C}$$

(b) Local Markov property:

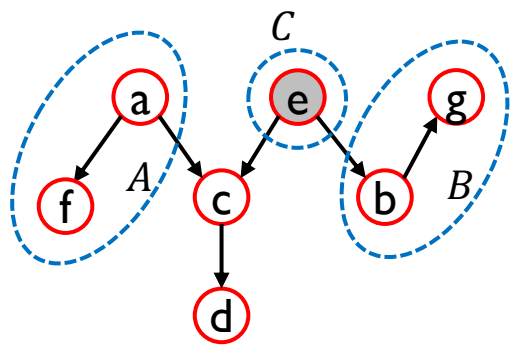
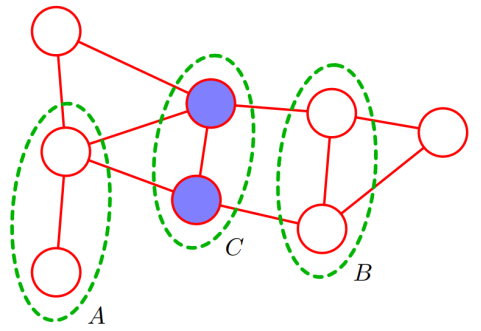
$$\mathbf{X}_v \perp \mathbf{X}_{V \setminus N[v]} \mid \mathbf{X}_{N(v)} \text{ for } \forall v \in V$$

(c) Pairwise Markov property:

$$\mathbf{X}_i \perp \mathbf{X}_j \mid \mathbf{X}_{V \setminus \{i, j\}} \text{ for } \forall (i, j) \notin E$$

- **Theorem:** the above three Markov properties are equivalent for a positive probability distribution.

Summary for Graphical Modeling

	Factorization Properties	Independent Properties
Bayesian Network	$p(\mathbf{x}) = \prod_{s=1}^N p(x_s \mathbf{x}_{\Gamma(s)})$ <p>where $\Gamma(s)$ denotes the set of parents of x_s.</p>	<p>D-separation</p> 
Markov Random Field	$p(\mathbf{x}) = \frac{1}{Z} \prod_c \phi_c(\mathbf{x}_c)$ <p>$\phi_c(\mathbf{x}_c)$ are potential functions of cliques.</p>	<p>Markov Properties</p> 

Outline

- Probabilistic Reasoning
- Graphical Models
 - Bayesian Networks
 - Markov Random Fields
- Conditional Independence
 - D-separation
 - Markov Properties
- **Inference in Graphical Models**
 - **Variable Elimination**
 - Belief Propagation

Probabilistic Inference

- Marginal inference

- What is the probability of a given variable in our model after we sum everything else out?

$$p(y = 1) = \sum_{x_1} \sum_{x_2} \cdots \sum_{x_n} p(y = 1, x_1, x_2, \dots, x_n)$$

- Maximum a posteriori (MAP) inference


- What is the most likely assignment to the variables in the model?

$$\max_{x_1, \dots, x_n} p(y = 1, x_1, \dots, x_n)$$

- It is **NP-hard** to answer these questions exactly for general cases.

Example: Markov Chain

- **Discrete assumption:** x_1, \dots, x_n are discrete taking k possible values.
- Given a chain Bayesian network (**Markov chain**).



- Interested in the marginal probability $p(x_n)$.

$$O(k^{n-1})!$$

$$p(x_n) = \sum_{x_1} \cdots \sum_{x_{n-1}} p(x_1, \dots, x_n) = \sum_{x_1} \cdots \sum_{x_{n-1}} p(x_1) \prod_{i=2}^n p(x_i | x_{i-1})$$

- Variable elimination by leveraging the **factorization property**.

$$\sum_{x_{n-1}} p(x_n | x_{n-1}) \sum_{x_{n-2}} p(x_{n-1} | x_{n-2}) \cdots \sum_{x_1} p(x_2 | x_1) p(x_1)$$

Depend on x_{n-1}, x_n
 x_1 is eliminated

Example: Markov Chain

- Define **factors** as follows:

$$\sum_{x_{n-1}} p(x_n | x_{n-1}) \underbrace{\sum_{x_{n-2}} p(x_{n-1} | x_{n-2}) \cdots}_{\tau(x_{n-1})} \underbrace{\sum_{x_1} p(x_2 | x_1) p(x_1)}_{\tau(x_2)}$$

1. **function** MARGINAL-INFERENCE(*Markov chain*) **returns** $p(x_n)$
2. **for** each assignment s to x_1 **do**
3. $\tau(x_1 = s) = p(x_1 = s)$ (1). Initialization
4. **for** t from 2 to n **do** $O(nk^2)$
5. **for** each assignment s to x_k **do**
6. $\tau(x_t = s) = \sum_{x_{t-1}} p(x_t = s | x_{t-1}) \tau(x_{t-1})$ (2). Elimination
7. **return** $p(x_n) = \tau(x_n)$.

Factors

- We are given a graphical model as a **product of factors**:

$$p(x_1, \dots, x_n) = \prod_{s \in F} \phi_s(\mathbf{x}_s).$$

Bayesian Networks	Markov Random Field
$p(\mathbf{x}) = \prod_{s=1}^N p(x_s \mathbf{x}_{\Gamma(s)})$ <p>$\Gamma(s)$ denotes parents of x_s.</p>	$p(\mathbf{x}) = \frac{1}{Z} \prod_c \phi_c(\mathbf{x}_c)$ <p>$\phi_c(\mathbf{x}_c)$: functions of cliques</p>
$F = \{s \cup \Gamma(s) \mid \forall s\}$ $\phi_s = p(x_s \mid \mathbf{x}_{\Gamma(s)})$	<p>F is the set of cliques</p> $\phi_s = \phi_c(\mathbf{x}_c) / Z'$

Operations

- Product

$$\phi_3(x_c) = \phi_1(x_c^{(1)}) \times \phi_2(x_c^{(2)})$$

- $x_c^{(i)}$ denotes an **assignment** to the variables in the scope of ϕ_i .
- For example, $\phi_3(a, b, c) = \phi_1(a, b) \times \phi_2(b, c)$.


- Marginalization

- If we have a factor $\phi(X, Y)$ over two sets of variables X, Y , marginalizing Y produces a new **factor**:

$$\tau(x) = \sum_y \phi(x, y)$$

Ordering

- In the example of **Markov chain**, we eliminate variables with the positive sequence ordering $(x_1 \rightarrow x_2 \rightarrow \dots \rightarrow x_{n-1})$.
- Important notations:
 - Different **orderings** may dramatically alter the running time of the variable elimination algorithm.



- It is NP-hard to find the **best ordering**.

Variable Elimination

1. **function** VARIABLE-ELIMINATION(*graphical model*) **returns** $p(x_n)$
2. **Initialize** factors set Φ of the graphical model.
3. Select an ordering O among variables except x_n . (1). Initialization
4. **for** each variable x_i ordered according to O **do**
5. $S = \{\phi_j \mid \phi_j \in \Phi, \phi_j \text{ depends on } x_i\}$ (2). Find relevant factors
6. $p_i = \prod_{\phi_j \in S} \phi_j$
7. $\tau = \sum_{x_i} p_i$ (3). Marginalization
8. $\Phi = \Phi \cup \{\tau\} \setminus S$
9. **return** $p(x_n) = \tau(x_n)$.


- Available for both directed (Bayesian) and undirected (Markov) graphs.

Variable Elimination

- Algorithm (Variable Elimination, VE)
 - For each variable X_i (ordered according to O):
 1. Multiply all factors ϕ_i containing X_i ;
 2. Marginalize out X_i to obtain a new factor τ ;
 3. Replace the factors ϕ_i with τ .

- Example: Alarm network

$$-\tau(x_2, x_3) = \sum_{x_1} p(x_2 \mid x_1, x_3) p(x_1)$$




Variable Elimination

- Algorithm (Variable Elimination, VE)
 - For each variable X_i (ordered according to O):
 1. Multiply all factors ϕ_i containing X_i ;
 2. Marginalize out X_i to obtain a new factor τ ;
 3. Replace the factors ϕ_i with τ .

- Example: Alarm network

$$- \tau(x_2, x_3) = \sum_{x_1} p(x_2 \mid x_1, x_3) p(x_1)$$

$$- \tau(x_2, x_4) = \sum_{x_3} \tau(x_2, x_3) p(x_3) p(x_4 \mid x_3)$$




Variable Elimination

- Algorithm (Variable Elimination, VE)
 - For each variable X_i (ordered according to O):
 1. Multiply all factors ϕ_i containing X_i ;
 2. Marginalize out X_i to obtain a new factor τ ;
 3. Replace the factors ϕ_i with τ .

- Example: Alarm network

- $\tau(x_2, x_3) = \sum_{x_1} p(x_2 | x_1, x_3) p(x_1)$
- $\tau(x_2, x_4) = \sum_{x_3} \tau(x_2, x_3) p(x_3) p(x_4 | x_3)$
- $p(x_2) = \tau(x_2) = \sum_{x_4} \tau(x_2, x_4)$



Introducing Evidence

- Consider a **general** distribution $P(X, Y, E)$ over sets of:
 - **Query** variables Y ;
 - **Observed** evidence variables E ;
 - **Unobserved** variables X .

$$P(Y \mid E = e) = \frac{P(Y, E = e)}{P(E = e)} \quad \left. \vphantom{\frac{P(Y, E = e)}{P(E = e)}} \right\} \begin{array}{l} \text{Apply VE} \\ \text{Algorithm!} \end{array}$$

- We can select the elimination ordering $X \rightarrow Y \rightarrow E$ to obtain $P(Y, E = e)$ and $P(E = e)$ in a single run of VE algorithm.
- How to **reutilize** the computation of $P(Y_1 \mid E_1 = e_1)$ for new query?
 - Example query: $P(Y_2 \mid E_2 = e_2)$.

Outline

- Probabilistic Reasoning
- Graphical Models
 - Bayesian Networks
 - Markov Random Fields
- Conditional Independence
 - D-separation
 - Markov Properties
- **Inference in Graphical Models**
 - Variable Elimination
 - **Belief Propagation**




Variable Elimination on Tree

- We can find the **optimal ordering** by ranking the nodes in **post-order**.
 - A post-order traversal visits a node after its children.
- At each step, we will eliminate a leaf node x_j .

$$-\tau_k(x_k) = \sum_{x_j} \phi(x_k, x_j) \tau_j(x_j)$$

Parent of x_j

summarizes all of the information from the subtree rooted at x_j




- $\tau_k(x_k)$ can be understood as the **message** delivered from x_j to x_k , denoted as $m_{j \rightarrow k}(x_k)$



Variable Elimination on Tree

- We want to compute $p(x_2)$ after $p(x_1)$:



- Variable elimination processes **share messages**.
- Given a graph $G = (V, E)$, if we store $m_{i \rightarrow j}(x_j)$ and $m_{j \rightarrow i}(x_i)$ for each $(i, j) \in E$, we can compute $p(x_i)$ with $O(1)$ steps in average.

Sum-Product Message Passing




- The message transmits from x_i to x_j :

$$m_{i \rightarrow j}(x_j) = \sum_{x_i} \phi(x_i) \phi(x_i, x_j) \prod_{l \in N(i) \setminus j} m_{l \rightarrow i}(x_i)$$

where $N(i) \setminus j$ refers to the set of nodes that are neighbors of i , excluding j .

- After we have computed all messages, we may answer any **marginal query** over x_i .



$$p(x_i) \propto \phi(x_i) \prod_{l \in N(i)} m_{l \rightarrow i}(x_i)$$



Message Passing Algorithm



1. **function** MESSAGE-PASSING(*tree model*) **returns** marginal functions
2. Pick an arbitrary node x_r as root.
3. Sort the nodes in **post-order** O .
-
4. **for** each non-root variable x_i ordered according to O **do**
5. Compute and propagate messages from x_i to its parent x_j . (a). Forward
-
6. **for** each non-root variable x_s ordered according to **reverse** O **do** (b). Backward
7. Compute and propagate messages from x_s to each of its children x_t
-
8. Compute the product of received messages at each node for x_n .



Message Passing for MAP Inference



- **Distributive law:** Sum and max operators distribute over products.
- Message passing can be used to perform MAP inference.
- Recall the example of Markov chain:

Marginal

$$p(x_n) = \sum_{x_1} \cdots \sum_{x_{n-1}} p(x_1) \prod_{i=2}^n p(x_i | x_{i-1})$$

Σ

$$\sum_{x_{n-1}} p(x_n | x_{n-1}) \sum_{x_{n-2}} p(x_{n-1} | x_{n-2}) \cdots \sum_{x_1} p(x_2 | x_1) p(x_1)$$



MAP

$$\max_{x_1, \dots, x_n} p(x_1, \dots, x_n) = \max_{x_1} \cdots \max_{x_n} p(x_1) \prod_{i=2}^n p(x_i | x_{i-1})$$

max

$$\max_{x_n} \max_{x_{n-1}} p(x_n | x_{n-1}) \cdots \max_{x_1} p(x_2 | x_1) p(x_1)$$



Max-Product Message Passing


- The message transmit from x_i to x_j :

$$m_{i \rightarrow j}(x_j) = \max_{x_i} \phi(x_i) \phi(x_i, x_j) \prod_{l \in N(i) \setminus j} m_{l \rightarrow i}(x_i)$$

where $N(i) \setminus j$ refers to the set of nodes that are neighbors of i , excluding j .

- After we have computed all messages, we may answer the MAP inference.

$$p^* = \max_{x_i} \phi(x_i) \prod_{l \in N(i)} m_{l \rightarrow i}(x_i)$$





Loopy Belief Propagation

- Inference on general graphs (**non-tree structure**):
 - **Approximate** solution (other methods are in next lectures)

1. **function** MESSAGE-PASSING(*graphical model*) **returns** messages
2. **Initialize** messages for each edge as uniform distribution.
3. Select an ordering for edges O .
4. **loop** until a fixed number of steps T or convergence **do**
5. $t \leftarrow$ the number of iterations
6. **for** each edge (i, j) ordered according to O **do**
7. $m_{i \rightarrow j}^{t+1}(x_j) = \sum_{x_i} \phi(x_i) \phi(x_i, x_j) \prod_{l \in N(i) \setminus j} m_{l \rightarrow i}^t(x_i)$
8. $m_{j \rightarrow i}^{t+1}(x_i) = \sum_{x_j} \phi(x_j) \phi(x_i, x_j) \prod_{l \in N(j) \setminus i} m_{l \rightarrow j}^t(x_j)$
9. **return** messages $m_{i \rightarrow j}^T(x_j), m_{j \rightarrow i}^T(x_i)$ for each edge (i, j)

Thank You

Questions?

Mingsheng Long

mingsheng@tsinghua.edu.cn

<http://ise.thss.tsinghua.edu.cn/~mlong>

答疑：东主楼11区413室